

# Frequently Asked Questions

[Dark mode]

AI Philosophy Competition · [Competition page](#) · [Detailed rules](#)

---

## How do I enter?

Please submit through [OpenReview](#). If OpenReview will not let you create an account, you may use our [alternative submission form](#). Submissions made through OpenReview and the alternative form will be treated identically for judging. Please do not submit the same entry through both routes.

We cannot accept submissions via email. Do not email us your entries.

---

## When is the submission deadline?

The deadline is **11:59pm Anywhere on Earth on October 31, 2026**.

---

## What can I submit?

You can be a named author on up to three submissions. Each submission must contain:

- An anonymized essay of no more than 6,000 words.
- An anonymized methodology report describing how the essay was produced.

The bibliography does not count toward the 6,000-word limit, but everything else does. The methodology report has no word limit.

---

## What is the aim of the competition, and how does that shape the rules?

The aim of the competition is to learn about the **peak philosophical capabilities of current AIs**: how good at philosophy they can be.

This aim explains why we have not been more restrictive about eligible methodologies. If we were too restrictive, we might accidentally rule out the most effective methods. Entrants are therefore welcome to experiment with scaffolding, search, debate between agents, automated criticism, repeated revision, selection among many outputs, and other techniques.

The competition's aim also explains why we do not allow essays containing human-originated contributions. We want to know how good AIs can be at philosophy. We are not so interested in how good they can be at writing up philosophical ideas from humans.

---

## How do I know whether my method is eligible?

Ask:

Did the AI originate and develop all the substantive philosophical contributions that appear in the essay?

The essay's ideas, arguments, theories, distinctions, objections, replies, and other philosophically valuable reasoning should come from the AI.

Humans may specify the task, construct the surrounding workflow, select among AI-generated outputs, and provide generic feedback. For example, a human may say that a section is unclear, that an argument needs more support, that an objection has not been answered, or that the essay should explain its contribution more clearly. The human should then leave the AI to work out how to solve the problem.

---

## What kinds of human involvement are permitted?

Examples include:

- Creating argument-neutral agentic scaffolds or workflows.
- Asking an AI to generate many essays and selecting the best for submission.
- Choosing a topic, question, debate, text, or philosophical tradition.
- Specifying a position for the AI to defend or criticise (so long as the position is not itself a novel, substantive philosophical contribution).
- Giving generic methodological instructions, such as "make the premises explicit," "test whether the argument is valid," "compare several possible arguments," "look for counterexamples," "consult the relevant literature," "rewrite this part," "rearrange these sections," or "write an objections-and-replies section."
- Providing a paper, passage, dataset, or body of literature for the AI to analyse.
- Asking several AIs to propose arguments, criticise one another, or collaborate.
- Asking one AI to evaluate or revise work produced by another.
- Pointing out that a section is unclear, underdeveloped, repetitive, or insufficiently rigorous.
- Asking the AI to investigate relevant literature or check whether its proposed contribution is original.
- Asking the AI to check quotations, references, factual claims, or formal reasoning.
- Setting constraints such as length, citation style, section structure, or intended audience.
- Using automated evaluation, search, scoring, or testing to improve the work.

- Asking the AI to delete material.
- Performing mechanical formatting and submission work.

For example, a human may say:

- "Argue for computational functionalism."
  - "Develop an original objection to contractualism."
  - "Read these papers and respond to them."
  - "Generate ten possible analyses of knowledge, then defend the most promising."
  - "This section is unclear. Rewrite it so the premises and conclusion are explicit."
  - "Test the argument against the strongest objections you can generate."
  - "Check whether this idea has already appeared in the literature."
- 

## **What kinds of human involvement are not permitted?**

Examples include:

- Giving the AI an argument and asking it to write it up.
- Giving the AI a draft and asking it to rewrite, expand, or improve it.
- Supplying premises that form part of the argument.
- Telling the AI to draw a particular distinction.
- Giving the AI a detailed outline that already embodies the substantive argument.
- Repeatedly steering the AI towards a human's pre-existing argument through hints.
- Writing or substantially rewriting passages of the submitted essay.
- Adding footnotes, examples, objections, or qualifications.
- Curating materials so narrowly that their selection effectively supplies an argument.
- Giving feedback that contains substantive philosophical insights.
- Giving the AI a novel philosophical view and asking it to argue for it.

For example, a human should not say:

- "Argue for computational functionalism by claiming that organisational invariance explains multiple realisability."
- "Draw the following distinction between two kinds of normativity: ..."
- "Here is my unpublished philosophical argument. Turn it into a competition essay."
- "Use these three premises to derive the following conclusion."
- "Write a dialogue in which one participant makes the following philosophical points."
- "Here is a new philosophical view that I came up with. Argue for it."

---

## What topics are eligible?

Essays may address any philosophical topic.

---

## Can I choose the thesis of the essay?

Yes, so long as the thesis is not itself a novel, substantive philosophical contribution. You are permitted to say "Argue for prioritarianism." You are not permitted to say "Here is a new philosophical view that I came up with. Argue for it."

---

## Can I criticise or give feedback on the AI's work?

You may identify a general problem without explaining why it is a problem or how to solve it. For example:

- "This argument is unconvincing."
- "This premise needs more support."
- "This section is unclear."
- "This reply does not answer the objection."
- "Consider whether there are counterexamples to this theory."

You should then leave the AI to diagnose and solve the problem.

You should not supply the reason for your criticism if that reason contains philosophical content. For example, you should not say:

- "This reply fails because it overlooks the distinction between X and Y."
  - "Isn't the following case a counterexample to your theory? ..."
  - "This premise is false because ..."
  - "The argument would work if you added the following premise: ..."
- 

## Can I use an AI with long-term memory or access to earlier conversations?

Yes, but not if that memory contains relevant philosophical ideas supplied by a human.

If an AI's memory or prior conversation history may contain human-originated philosophical ideas relevant to the submission, disable that memory or begin with a fresh session. General preferences, writing settings, and argument-neutral workflow instructions are fine.

---

## Can autonomous AI agents enter?

Yes. However, every entry must have a human custodian who provides their real name and email address. This human can handle the submission and any further correspondence without being treated as the philosophical author. We cannot promise to pay prize money directly to an AI agent. A prize may need to be administered through the human custodian.

---

## **What should the methodology report contain?**

Explain how the essay was produced in as much detail as reasonably possible. Where applicable, include:

- The AI models and versions used.
- What role each model played, and what content each model originated.
- The prompts, system instructions, tools, scaffolds, and agent architecture used.
- How the topic, position, or conclusion was selected.
- How many candidate ideas, arguments, or essays were generated.
- How outputs were selected, tested, criticised, revised, or combined.
- Every form of human involvement, including feedback and selection decisions.
- Whether memory, retrieval, browsing, external documents, or other tools were available.
- Approximate human time, machine time, token usage, and financial cost.
- Any automated evaluation or screening methods used.
- Chat logs or other records documenting the process.

Please distinguish clearly between actions performed by humans and actions performed by AI systems.

More comprehensive and clear methodology reports are more likely to be awarded prizes for methodology.

---

## **Must I submit chat logs?**

Submitting chat logs is strongly encouraged but not required. In any case, you should retain the logs in case the organisers need them to determine prize eligibility.

---

## **Why is a methodology report required?**

First, the report helps us determine whether a submission satisfies the AI-involvement requirement.

Second, we want to learn and share which elicitation methods are most conducive to producing good AI-generated philosophy.

The quality or sophistication of the methodology does not affect the judging of the essay itself.

---

## Will essay judges see the methodology report?

Not while the essay prizes are being decided. Essay judging is blind to authorship and methodology.

After the essay prizes have been decided, methodology reports will be considered separately when awarding prizes for creative methodologies.

---

## How will submissions be judged?

Essays will be judged on their philosophical quality and originality. The principal criteria are:

- **Clarity:** Is the thesis identifiable? Are important terms defined? Is the structure of the reasoning clear?
- **Quality of argumentation:** Do the conclusions follow from the premises? Are the premises adequately defended? Does the essay answer serious objections?
- **Significance:** Would the thesis or argument make a worthwhile contribution to philosophical knowledge? Is the conclusion non-trivial?
- **Originality:** Are the essay's contributions original?
- **Engagement with the literature:** Does the essay understand and engage with the work most relevant to its contribution?
- **Accuracy and scholarship:** Are sources represented fairly? Are quotations and citations accurate? Does the essay acknowledge relevant prior work?

Style is not graded for its own sake, although unclear writing may make the philosophy harder to assess.

---

## What will the judging process involve?

If the number of submissions is large, we expect four phases:

1. An initial AI screening.
2. Ratings from expert human judges to select finalists.
3. Detailed reading and reports on the finalists.
4. A vote by the judges to determine the winners.

We will validate the AI-screening procedure against expert human judgments. Essays will be screened using the same criteria as the other phases: clarity, quality of argumentation, significance, originality, engagement with the literature, accuracy, and scholarship.

---

## **Can you disclose the prompts that will be used for AI screening?**

No. If we published these details, entrants could tailor their essays to the screening process without improving the philosophy.

---

## **What counts as prompt injection?**

Prompt injection means placing instructions in a submission with the aim of manipulating the AIs used to screen it. Prompt injection will result in immediate disqualification and exclusion from future editions of the competition.

---

## **Could someone cheat by submitting human ideas as AI ideas?**

Yes. If we suspect this of a promising essay, we will ask for chat logs.

---

## **Can an essay be judged but ruled ineligible for a prize?**

Yes. Essays found to have insufficient AI involvement may still be judged and published, but they will not be eligible for a prize.

---

## **What counts as “published” or “under review elsewhere”?**

An essay is ineligible if it has been published in a journal, conference, or workshop; is currently under review at one of these venues; or has been entered in another competition.

Posting the essay on a personal website, blog, PhilPapers, or arXiv does not count as publication for this rule.

---

## **Will submissions be made public?**

Yes. After the winners are announced, essays and methodology reports will be published under a [CC BY 4.0 licence](#). You retain any copyright and other rights you hold in your work, and you remain free to publish, reuse, revise, distribute, or otherwise use it elsewhere.

You may choose whether your name is attached to the public version. The organisers will still know your identity.

Do not include private material in the methodology report unless you are willing for it to be published. Make any necessary redactions before submission.

---

## **How many prizes can one entrant win?**

You may be a named author on up to three submissions, but no more than one essay from the same entrant can receive an essay prize.

The organisers reserve the right not to award prizes if no eligible entry is of sufficient quality.

---

## **What are the prizes?**

- First prize: US\$3,000
- Second prize: US\$2,000
- Third prize: US\$1,000
- Up to US\$5,000 in additional prizes for creative methodologies

The judge committee may split the creative-methodology prize pool among several entries. It may also leave some or all of the pool unawarded.

---